

# ParaToric 1.0-beta: Continuous-time quantum Monte Carlo for the toric code in a parallel field

Simon M. Linsel<sup>\*</sup> and Lode Pollet<sup>†</sup>

Department of Physics and Arnold Sommerfeld Center for Theoretical Physics (ASC),  
Ludwig-Maximilians-Universität München, Theresienstr. 37, München D-80333, Germany  
Munich Center for Quantum Science and Technology (MCQST), Schellingstr. 4, D-80799  
München, Germany

<sup>\*</sup> [simon.linsel@lmu.de](mailto:simon.linsel@lmu.de), <sup>†</sup> [lode.pollet@lmu.de](mailto:lode.pollet@lmu.de)

## Abstract

We introduce ParaToric, a C++ package for simulating the toric code in a parallel field (i.e., X- and Z-fields) at finite temperature. We implement and extend the continuous-time quantum Monte Carlo algorithm of Wu, Deng, and Prokof'ev on the square, triangular, honeycomb, and cubic lattices with open and periodic boundaries, respectively. The package is expandable to arbitrary lattice geometries and custom observables diagonal in either the X- or Z-basis. ParaToric also supports snapshot extraction in both bases, making it ideal for generating training/benchmarking data for other methods, such as lattice gauge theories, cold atom or other quantum simulators, quantum spin liquids, artificial intelligence, and quantum error correction. The software provides bindings to C/C++ and Python, and is thus almost universally integrable into other software projects.

Copyright attribution to authors.

This work is a submission to SciPost Physics Codebases.

License information to appear upon publication.

Publication information to appear upon publication.

Received Date

Accepted Date

Published Date

1

## Contents

1	<b>Introduction</b>	3
2	<b>The toric code in a parallel field</b>	3
3	2.1 Hamiltonian	3
4	2.2 Lattice geometries	3
5	2.3 Observables	4
6	<b>Installation &amp; interfaces</b>	6
7	3.1 C++ interface	6
8	3.1.1 Build & Installation	7
9	3.1.2 Public class <code>ExtendedToricCode</code>	8
10	3.1.3 Configuration type	8
11	3.1.4 Return type	11
12	3.1.5 C++ usage examples	11
13	3.2 C++ command-line interface	13
14	3.2.1 Build & Installation	13

17	3.2.2	etc_sample	15
18	3.2.3	etc_hysteresis	16
19	3.2.4	etc_thermalization	16
20	3.2.5	HDF5 structure	16
21	3.3	C interface	16
22	3.3.1	Build & Installation	16
23	3.3.2	Status & error handling	17
24	3.3.3	Opaque handle	17
25	3.3.4	Configuration type	17
26	3.3.5	Return type	18
27	3.3.6	Procedures (mirror the C++ API)	19
28	3.3.7	C usage example	19
29	3.4	Python bindings	21
30	3.4.1	Build & Installation	21
31	3.4.2	Module layout	22
32	3.4.3	Usage example	24
33	3.5	Python command-line interface	24
34	3.5.1	Build & Installation	25
35	3.5.2	$T$ -sweep	26
36	3.5.3	$h$ -sweep	27
37	3.5.4	$\lambda$ -sweep	27
38	3.5.5	$\phi$ -sweep	28
39	3.5.6	Hysteresis-sweep	28
40	3.5.7	Thermalization	29
41	4	Using ParaToric	29
42	4.1	Monte Carlo Updates	29
43	4.2	Monte Carlo Diagnostics	30
44	4.2.1	Thermalization mode	30
45	4.2.2	Integrated autocorrelation time	30
46	4.2.3	Error bars	31
47	4.3	Tips & tricks	31
48	4.3.1	Probing ground state physics	31
49	4.3.2	Probing first-order transitions	31
50	4.3.3	Choosing the basis	31
51	4.3.4	Choosing $N_{\text{thermalization}}$	31
52	4.3.5	Choosing $N_{\text{samples}}$	31
53	4.3.6	Choosing $N_{\text{between\_samples}}$	32
54	4.3.7	Choosing $N_{\text{resamples}}$	32
55	4.3.8	Extracting snapshots	32
56	4.3.9	Adding new observables/lattices/updates	32
57	4.4	Benchmarks	32
58	4.4.1	Thermalization	32
59	4.4.2	Integrated autocorrelation time	33
60	4.4.3	Run-time	33
61	4.4.4	Topological phase transition	35
62	5	Conclusion & Outlook	36
63		References	37

## 1 Introduction

The toric code is one of the most fundamental and most-studied models in modern condensed matter physics. It was first written down by Kitaev [1] and is the simplest example of a model hosting a topological phase (a gapped  $\mathbb{Z}_2$  quantum spin liquid) and anyonic excitations. The toric code is also the foundational model for error-correcting codes [2, 3] and has deep connections to the Ising gauge theory [4].

The toric code can be extended with fields which, when strong enough, destroy the topological order. This model is sign-problem-free, thus making quantum Monte Carlo the method of choice. Wu, Deng, and Prokof'ev developed a continuous-time quantum Monte Carlo algorithm [5]. ParaToric implements and extends this algorithm with new updates which enable ergodicity at large temperatures and at zero off-diagonal field, thus significantly improving the applicability of the algorithm.

ParaToric implements a wide range of lattices, boundary conditions, and observables. It is also possible to extend ParaToric with new interactions, observables, and lattices. We provide documented interfaces in C, C++, and Python as well as command-line interfaces, making the integration of ParaToric into other projects and programming languages straightforward. ParaToric will save simulation results to HDF5 files and snapshots to GraphML files (XML-based), with a focus on interoperability with other packages. ParaToric comes with an MIT license.

## 2 The toric code in a parallel field

### 2.1 Hamiltonian

ParaToric implements and extends the continuous-time quantum Monte Carlo (QMC) algorithm by Wu, Deng, and Prokof'ev [5] to simulate the toric code in a parallel field (also called perturbed toric code or extended toric code)

$$\hat{\mathcal{H}} = -\mu \sum_v \hat{A}_v - J \sum_p \hat{B}_p - h \sum_l \hat{\sigma}_l^x - \lambda \sum_l \hat{\sigma}_l^z, \quad (1)$$

where  $J, \lambda > 0$  in the  $\hat{\sigma}^x$ -basis and  $\mu, h > 0$  in the  $\hat{\sigma}^z$ -basis (otherwise the model has a sign-problem).  $\hat{\sigma}_l^x$  and  $\hat{\sigma}_l^z$  are Pauli matrices defined on the links of the underlying lattice. The star term  $\hat{A}_v$  contains all links adjacent to lattice site  $v$ , the plaquette term  $\hat{B}_p$  contains all links that belong to the same elementary plaquette  $p$  of the underlying lattice. The temperature  $T = 1/\beta$  is finite. For readers interested in extending the code, we note that it is relatively straightforward to add interactions that are diagonal in the chosen basis, such as (long-range) Ising interactions. Off-diagonal interactions require a more careful review and extension of the Monte Carlo updates to ensure ergodicity. However, diagonal interactions can also lead to sampling problems, especially when they introduce frustration.

### 2.2 Lattice geometries

We implement the square, honeycomb, triangular, and cubic lattices, see Fig. 1. On the cubic lattice, the plaquettes contain the four links of cube faces, *not* the twelve links of the cube




Figure 1: **Implemented lattices.** We implement the extended toric code (1) on the square (a), honeycomb (b), triangular (c), and cubic lattices (d). For each lattice, we show the star ( $\hat{A}_v$ ) and plaquette ( $\hat{B}_p$ ) terms. The cubic lattice importantly features star interactions of six links and plaquette interactions of four links on the faces of cubes.

102 (that model has a different m-anyon structure). We implement open and periodic boundaries,  
 103 respectively. New lattices can be added in `src/lattice/lattice.cpp`.

### 104 2.3 Observables

105 Here we list all observables that ParaToric implements for the extended toric code. Custom  
 106 observables can be added in `src/mcmc/extended_toric_code_qmc.hpp`. For each ob-  
 107 servable  $\hat{O}$ , we calculate the expectation value  $\langle \hat{O} \rangle$  and the Binder ratio  $U_O = \frac{\langle \hat{O}^4 \rangle}{\langle \hat{O}^2 \rangle^2}$  with error  
 108 bars obtained from bootstrapping (see below), respectively.

109 `anyon_count` The number of e-anyons ( $\hat{\sigma}^x$ -basis) or m-anyons ( $\hat{\sigma}^z$ -basis) in the system.

110 `anyon_density`  $\hat{\sigma}^x$ -basis: The number of e-anyons divided by the number of lattice sites.  
 111  $\hat{\sigma}^z$ -basis: The number of m-anyons divided by the number of plaquettes.

112 `delta` The difference of the star and plaquette expectation values:  $\Delta = \langle \hat{A}_v \rangle - \langle \hat{B}_p \rangle$ .

113 `energy` The total energy  $E = \langle \hat{\mathcal{H}} \rangle$ .

114 `energy_h` The electric field term  $E_h = \langle -h \sum_l \hat{\sigma}_l^x \rangle$ .

115 `energy_lambda` The gauge field term  $E_\lambda = \langle -\lambda \sum_l \hat{\sigma}_l^z \rangle$ . We write  $\lambda$  as `lambda` because some  
 116 programming languages feature `lambda` as a keyword.

117 `energy_J` The plaquette term  $E_J = \langle -J \sum_p \hat{B}_p \rangle$ .

118 `energy_mu` The star term  $E_\mu = \langle -\mu \sum_v \hat{A}_v \rangle$ .

119 `fredenhagen_marcu` The equal-time Fredenhagen-Marcu loop operator [6–8]:

$$O_{\text{FM}}^{x/z} = \lim_{L \rightarrow \infty} \frac{\langle \prod_{l \in C_{1/2}^{x/z}} \hat{\sigma}_l^{x/z} \rangle}{\sqrt{|\langle \prod_{l \in C^{x/z}} \hat{\sigma}_l^{x/z} \rangle|}}, \quad (2)$$

120  $C_{1/2}^{x/z}$  is half,  $C^{x/z}$  is a full Wilson loop in the  $\hat{\sigma}^x$ -basis ('t Hooft loop in the  $\hat{\sigma}^z$ -basis). The loop is  
 121 automatically constructed for all supported lattices, and the perimeter scales with  $\mathcal{O}(L)$ , where  
 122  $L$  is the linear system size. When probing perimeter/area laws, the user should change  $L$ . We  
 123 currently do not support off-diagonal loop operators, e.g., measuring products of  $\hat{\sigma}^x$ -operators  
 124 in the  $\hat{\sigma}^z$ -basis.

125 `largest_cluster` The largest connected cluster of neighboring bonds with  $\hat{\sigma}^x = -1$   
 126 ( $\hat{\sigma}^z = -1$ ) in the  $\hat{\sigma}^x$ -basis ( $\hat{\sigma}^z$ -basis). This observable is used to calculate the percolation  
 127 strength, see [9].

128 `percolation_probability` Measures the bond percolation probability, i.e. if we can  
 129 wind around the system while only traversing bonds with  $\hat{\sigma}^x = -1$  ( $\hat{\sigma}^z = -1$ ) in the  $\hat{\sigma}^x$ -basis  
 130 ( $\hat{\sigma}^z$ -basis). Formally, it is the expectation value  $\langle \hat{\Pi}^{x/z} \rangle$  of the projector

$$\hat{\Pi}^{x/z} = \sum_{W(j) \neq 0} |\{\hat{\sigma}^{x/z}\}_j\rangle \langle \{\hat{\sigma}^{x/z}\}_j|, \quad (3)$$

131 over all possible configurations  $\{\hat{\sigma}^{x/z}\}_j$  with non-zero winding number  $W(j)$  of connected link  
 132 clusters of neighboring  $\hat{\sigma}^{x/z} = -1$ . These clusters are called percolating clusters. For details,  
 133 see [9–11].

134 `percolation_strength` If a snapshot does not have a percolating cluster, the percolation  
 135 strength is 0. If a snapshot has a percolating cluster, the percolation strength is defined as the  
 136 result of `largest_cluster` divided by the total number of links in the system. For details,  
 137 see [9, 11].

138 `plaquette_percolation_probability` Similar to the percolation probability of bonds.  
 139 Two plaquettes are in the same cluster if they share a link  $l$  with  $\hat{\tau}_l^x = -1$ . For details, see [11].

140 `plaquette_z` The plaquette expectation value  $\langle \hat{B}_p \rangle$ .

141 `sigma_x` The electric field expectation value  $\langle \hat{\sigma}^x \rangle$ .

142 `sigma_x_susceptibility` The static susceptibility

$$\chi^x = \frac{1}{N} \int_0^\beta \langle \hat{\sigma}^x(0) \hat{\sigma}^x(\tau) \rangle_c d\tau, \quad (4)$$

143 where  $N$  is the total number of links (qubits) and the integral is over the imaginary time  $\tau$ <sup>1</sup>.  
 144 Importantly,  $\chi^x$  can be calculated both in the  $\hat{\sigma}^x$ - and the  $\hat{\sigma}^z$ -basis.

<sup>1</sup>The dynamical (fidelity) susceptibility is a trivial extension and contains an extra  $\tau$  dependency in the integral.

145 `sigma_z` The gauge field expectation value  $\langle \hat{\sigma}^z \rangle$ .

146 `sigma_z_susceptibility` The static susceptibility

$$\chi^z = \frac{1}{N} \int_0^\beta \langle \hat{\sigma}^z(0) \hat{\sigma}^z(\tau) \rangle_c d\tau, \quad (5)$$

147 where  $N$  is the total number of links (qubits) and the integral is over the imaginary time  $\tau$ .

148 Importantly,  $\chi^z$  can be calculated both in the  $\hat{\sigma}^x$ - and the  $\hat{\sigma}^z$ -basis.

149 `staggered_imaginary_times` Order parameter from [5]. It is defined as

$$O_{\text{SIT}}^{x/z} = \frac{1}{\beta} [(\tau_1^k - 0) - (\tau_2^k - \tau_1^k) + \dots + (-1)^{N(k)-1} (\tau_{N(k)}^k - \tau_{N(k)-1}^k) + (-1)^{N(k)} (\beta - \tau_{N(k)}^k)], \quad (6)$$

150 where  $\tau_n^k$  is the imaginary time of the  $n$ -th tuple spin flip of type  $k$ .  $k$  is a plaquette  $p$  (star  
151  $s$ ) of links in the  $\hat{\sigma}^x$ -basis ( $\hat{\sigma}^z$ -basis). This order parameter can neither be evaluated from  
152 snapshots, nor from any other method that does not have access to imaginary time.

153 `star_x` The star expectation value  $\langle \hat{A}_v \rangle$ .

154 `string_number` The total number of links with  $\hat{\sigma}^x = -1$  in the  $\hat{\sigma}^x$ -basis ( $\hat{\sigma}^z = -1$  in the  
155  $\hat{\sigma}^z$ -basis).

### 156 3 Installation & interfaces

157 There are five ways to use paratoric, directly from within code (C, C++, Python) or via the  
158 command-line (C++, Python). All interfaces require compiling C++ code. We tested the  
159 compilation with GCC 15 and Clang 20.

160 All interfaces implement three functionalities. *Thermalization* simulations are used to  
161 benchmark the thermalization process of a Markov chain and are primarily a diagnostic tool.  
162 Regular *sampling* routines are used for generating snapshots and measuring observables, e.g.,  
163 in the context of continuous phase transitions. *Hysteresis* routines are a variant of the regu-  
164 lar sampling routines where not one but an array of Hamiltonian parameters is provided and  
165 only one Markov chain is used for all parameters. The order of the Hamiltonian parameters  
166 in the input array matters: The last state of the previous parameter is used as an initial state  
167 for the thermalization phase of the next parameters. This simulation type should primarily be  
168 used when mapping out hysteresis curves in the vicinity of first-order phase transitions, hence  
169 the name. Since the hysteresis simulation returns the values of not one but many parameter  
170 sets, the output types are generally different from the regular sampling. It is also much slower  
171 than regular sampling, because the simulation for different parameters can in general not be  
172 parallelized.

#### 173 3.1 C++ interface

174 The C++ interface enables users to use a ParaToric public header from within another C++  
175 project.

### 176 3.1.1 Build & Installation

177 The core requires C++23, CMake  $\geq 3.23$ , and Boost  $\geq 1.87$  (older Boost versions may work,  
178 but were not tested). To compile it, run:

```
179 cmake -S . -B build -DCMAKE_BUILD_TYPE=Release \
180 -DPARATORIC_ENABLE_NATIVE_OPT=ON -DPARATORIC_LINK_MPI=OFF \
181 -DPARATORIC_BUILD_TESTS=ON
182 cmake --build build -jN
183 ctest --test-dir build -jN --output-on-failure
184 cmake --install build
```

185 Replace N with the number of cores to use, e.g. -j4 for 4 cores.

- 186 • -DCMAKE\_BUILD\_TYPE=Release. Only set to Debug if you're a developer.
- 187 • -DCMAKE\_INSTALL\_PREFIX. By default, executables install to  $\{\text{CMAKE\_SOURCE\_DI}$   
188  $\text{R}\}/\{\text{CMAKE\_INSTALL\_BINDIR}\}/$ , headers to  $\{\text{CMAKE\_INSTALL\_INCLUDEDIR}\}/\text{p}$   
189  $\text{aratoric}$ , and static libraries to  $\{\text{CMAKE\_SOURCE\_DIR}\}/\{\text{CMAKE\_INSTALL\_LIBD}$   
190  $\text{IR}\}/$ . The Python scripts expect  $\{\text{CMAKE\_SOURCE\_DIR}\}/\text{bin}/$ ; this directory always  
191 contains the paratoric executable. To install into a custom directory, pass it via -DCM  
192  $\text{AKE\_INSTALL\_PREFIX}$ , e.g. -DCMAKE\_INSTALL\_PREFIX=/your/custom/directo  
193  $\text{ry}/$ .
- 194 • -DPARATORIC\_EXPORT\_COMPILE\_COMMANDS=ON. Export compile\_commands.json  
195 for tooling.
- 196 • -DPARATORIC\_LINK\_MPI=OFF. Link the core to MPI, required on some clusters. The  
197 core itself does not need MPI.
- 198 • -DPARATORIC\_ENABLE\_NATIVE\_OPT=OFF. Turn on -march=native on GCC and  
199 Clang.
- 200 • -DPARATORIC\_ENABLE\_AVX2=OFF. Enable AVX2 (Haswell New Instructions). Requires  
201 a CPU which supports AVX2.
- 202 • -DPARATORIC\_BUILD\_TESTS=PROJECT\_IS\_TOP\_LEVEL. Compile the tests (recom-  
203 mended).

#### 204 CMake usage (installed package)

```
205 cmake_minimum_required(VERSION 3.23)
206 project(my_qmc_app CXX)
207
208 find_package(paratoric CONFIG REQUIRED)    # provides paratoric::core
209
210 add_executable(myapp main.cpp)
211 target_link_libraries(myapp PRIVATE paratoric::core)
```

#### 212 CMake usage (as subdirectory)

213 If the core lives in deps/paratoric, add it and link to the same target:

```
214 add_subdirectory(deps/paratoric)
215 add_executable(myapp main.cpp)
216 target_link_libraries(myapp PRIVATE paratoric::core)
```

### 217 3.1.2 Public class ExtendedToricCode

218 The interface class ExtendedToricCode lives in the public header `#include<paratoric`  
 219 `/mcmc/extended_toric_code.hpp>`. All symbols are in the `paratoric` namespace. All  
 220 methods are static, take a single `Config` object and return a `Result` object. The required  
 221 fields in `config` are documented for each method within the docstrings.

222 `Result ExtendedToricCode::get_thermalization(Config config)` Run ther-  
 223 malization only. Required fields: `lat_spec.{basis,lattice_type,system_size,beta`  
 224 `,boundaries,default_spin}`, `param_spec.{mu,h,J,lmbda,h_therm,lmbda_therm`  
 225 `}`, `sim_spec.{N_thermalization,N_resamples,custom_therm,observables,seed`  
 226 `}`, `out_spec.{path_out,save_snapshots}`.

227 `Result ExtendedToricCode::get_sample(Config config)` Run a production mea-  
 228 surement pass. Returns the observables selected in `config`. Required fields: `lat_spec.{ba`  
 229 `sis,lattice_type,system_size,beta,boundaries,default_spin}`, `param_spec.`  
 230 `{mu,h,J,lmbda}`, `sim_spec.{N_samples,N_thermalization,N_between_samples,`  
 231 `N_resamples,observables,seed}`, `out_spec.{path_out,save_snapshots}`.

232 `Result ExtendedToricCode::get_hysteresis(Config config)` Perform a hys-  
 233 teresis sweep, where the last state of the previous parameter is used as the initial state of the  
 234 following parameter in `h_hys` & `lmbda_hys`. Required fields: `lat_spec.{basis,lattice`  
 235 `_type,system_size,beta,boundaries,default_spin}`, `param_spec.{mu,J,h_hys`  
 236 `,lmbda_hys}`, `sim_spec.{N_samples,N_thermalization,N_between_samples,N_r`  
 237 `esamples,observables,seed}`, `out_spec.{paths_out,save_snapshots}`.

### 238 3.1.3 Configuration type

239 The struct `Config` (declared in `<paratoric/types/types.hpp>`) contains multiple nested  
 240 specifications.

#### 241 Top-level configuration: Config

Field	Type	Purpose
sim_spec	SimSpec	Simulation / MC controls (backend-consumed).
param_spec	ParamSpec	Model couplings / parameters (backend-consumed).
lat_spec	LatSpec	Lattice geometry and basis.
out_spec	OutSpec	Output folders and snapshot toggles.



243 **Simulation specification** (config.sim\_spec)

244

Field	Type	Meaning / Defaults
N_samples	int	Number of recorded snapshots. Default 1000.
N_thermalization	int	Number of warmup steps before sampling. Typically $O(L^d)$ , where $L$ is the system size and $d$ is the dimensionality. Default 10000.
N_between_samples	int	Steps between consecutive snapshots. Higher value decreases autocorrelation and improves error bars. Typically $O(L^d)$ , where $L$ is the system size and $d$ is the dimensionality. Default 1000.
N_resamples	int	Bootstrap resamples for errors. Default 1000.
custom_therm	bool	Use custom thermalization schedule. Default false.
seed	int	PRNG seed. 0 means “random seed.” Default 0.
observables	vector<string>	Names of observables to record each snapshot. For options, see Sec. 2.3.

246 **Parameter specification** (config.param\_spec)

247

Field	Type	Meaning / Defaults
mu	double	Star term coefficient. Default 1.0.
h	double	Electric field term. Default 0.0.
J	double	Plaquette term. Default 1.0.
lmbda	double	Gauge-field term. Default 0.0.
h_therm	double	Thermalization value for h when using custom schedules. Default NaN (unused).
lmbda_therm	double	Thermalization value for lmbda when using custom schedules. Default NaN (unused).
h_hys	vector<double>	Sweep values of h for hysteresis runs. Default empty. Length must match lmbda_hys.
lmbda_hys	vector<double>	Sweep values of lmbda for hysteresis runs. Default empty. Length must match h_hys.

248 **Lattice specification** (config.lat\_spec)

249

Field	Type	Meaning / Valid values
basis	char	Spin eigenbasis for the simulation. Must be 'x' or 'z'.
lattice_type	string	The lattice ("square", "triangular", "honeycomb" or "cubic").
system_size	int	Linear system size (per dimension).
beta	double	Inverse temperature $\beta > 0$ .
boundaries	string	Boundary condition: "periodic" or "open".
default_spin	int	Initial link spin, must be +1 or -1.

250 **Output specification** (`config.out_spec`)

251

Field	Type	Meaning
path_out	string	Primary output folder name.
paths_out	vector<string>	Hysteresis subfolder names. Length must match h_hys.
save_snapshots	bool	Save snapshots toggle. Default false.

252

## 253 3.1.4 Return type

Field	C++ Type	Meaning
series	<code>vector&lt;vector&lt;variant&lt;complex&lt;double&gt;,double&gt;&gt;&gt;</code>	Time series of all requested observables, thermalization is excluded (except for thermalization simulation). Outer index = observable, inner index = time point.
acc_ratio	<code>vector&lt;double&gt;</code>	Time series of Monte Carlo acceptance ratios.
mean	<code>vector&lt;double&gt;</code>	Bootstrap observable means.
mean_std	<code>vector&lt;double&gt;</code>	Bootstrap standard errors of the mean.
binder	<code>vector&lt;double&gt;</code>	Bootstrap binder ratios.
binder_std	<code>vector&lt;double&gt;</code>	Bootstrap standard errors of the binder ratios.
tau_int	<code>vector&lt;double&gt;</code>	Estimated integrated autocorrelation times.
series_hys	<code>vector&lt;vector&lt;vector&lt;variant&lt;complex&lt;double&gt;,double&gt;&gt;&gt;</code>	Hysteresis time series of all requested observables, thermalization is excluded (except for thermalization simulation). Outer vector = hysteresis parameters (order as in <code>h_hys</code> , <code>lambda_hys</code> ), middle vector = observables (order as in <code>observables</code> ), inner vector = time series.
mean_hys	<code>vector&lt;vector&lt;double&gt;&gt;</code>	Hysteresis bootstrap observable means. Outer vector = hysteresis parameters (order as in <code>h_hys</code> , <code>lambda_hys</code> ), inner vector = observables (order as in <code>observables</code> ).
mean_std_hys	<code>vector&lt;vector&lt;double&gt;&gt;</code>	Hysteresis bootstrap standard errors of the mean. Indices as above.
binder_hys	<code>vector&lt;vector&lt;double&gt;&gt;</code>	Hysteresis bootstrap binder ratios. Indices as above.
binder_std_hys	<code>vector&lt;vector&lt;double&gt;&gt;</code>	Hysteresis bootstrap standard errors of the binder ratios. Indices as above.
tau_int_hys	<code>vector&lt;vector&lt;double&gt;&gt;</code>	Hysteresis estimated integrated autocorrelation times. Indices as above.

## 255 3.1.5 C++ usage examples

Listing 1: C++ API - Minimal call

```

256 // C++23
257 #include <iostream>
258 #include <print>
259 #include <vector>
260 #include <string>
261 #include <paratoric/mcmc/extended_toric_code.hpp>
262 #include <paratoric/types/types.hpp>
263
264
265 int main() {
266     using namespace paratoric;
267
268     Config cfg{};
269
270     // ---- lattice sub-config (required) ----
271     cfg.lat_spec.basis = 'z'; // or 'x'
272     cfg.lat_spec.lattice_type = "square"; // or "cubic", "honeycomb", ...
273     cfg.lat_spec.system_size = 16;
274     cfg.lat_spec.beta = 8.0;
275     cfg.lat_spec.boundaries = "periodic"; // or "open"
276     cfg.lat_spec.default_spin = 1;
277
278     // ---- Hamiltonian parameters ----
279     cfg.param_spec.mu = 1.0; // star term
280     cfg.param_spec.J = 1.0; // plaquette term
281     cfg.param_spec.h = 0.20; // electric field term
282     cfg.param_spec.lmbda = 0.00; // gauge-field term
283
284     // Optional thermalization schedule values (used if custom_therm = true)
285     cfg.param_spec.h_therm = std::numeric_limits<double>::quiet_NaN();
286     cfg.param_spec.lmbda_therm = std::numeric_limits<double>::quiet_NaN();
287
288     // (Optional) Hysteresis sweep grids - only read by get_hysteresis(...)
289     cfg.param_spec.h_hys = {}; // e.g. {0.0, 0.1, 0.2, 0.3, 0.2, 0.1, 0.0}
290     cfg.param_spec.lmbda_hys = {}; // e.g. {0.0, 0.1, 0.2, 0.3, 0.2, 0.1, 0.0}
291
292     // ---- Simulation (MC) controls ----
293     cfg.sim_spec.N_samples = 0; // 0 => thermalization-only
294     cfg.sim_spec.N_thermalization = 5000; // warmup steps
295     cfg.sim_spec.N_between_samples = 10; // thinning between snapshots
296     cfg.sim_spec.N_resamples = 1000; // bootstrap
297     cfg.sim_spec.custom_therm = false; // set true to use *_therm values
298     cfg.sim_spec.seed = 12345; // 0 => random seed
299
300     // Observables to record each snapshot (backend-recognized names)
301     cfg.sim_spec.observables = {
302         "energy", // total energy
303         "plaquette_z", // plaquette energy
304         "anyon_count", // number of anyons (x-basis: e-anyons, z-basis: m-anyons)
305     }
306     "fredenhagen_marcu" // example: Wilson/'t Hooft loop proxy
307 };
308
309 // ---- Output / I/O policy ----
310 cfg.out_spec.path_out = "runs/sample"; // single-run output dir
311 cfg.out_spec.paths_out = {}; // filled only for hysteresis
312 cfg.out_spec.save_snapshots = false; // set true to dump every snapshot

```

```

313   cfg.out_spec.full_time_series = true; // save full time series (FCS)
314
315   // 1) Check thermalization
316   Result warmup = ExtendedToricCode::get_thermalization(cfg);
317   std::print("Thermalization_\u00series:\u00{}", warmup.series);
318
319   // 2) Production sample (set N_samples > 0 and call get_sample)
320   cfg.sim_spec.N_samples = 2000;
321   Result out = ExtendedToricCode::get_sample(cfg);
322   std::print("Production_\u00autocorrelations:\u00{}", out.tau_int);
323
324   return 0;
325 }

```

## 3.2 C++ command-line interface

ParaToric ships a C++ command-line interface `${CMAKE_INSTALL_PREFIX}/${CMAKE_INSTALL_BINDIR}/paratoric` that orchestrates C++ backends, runs sweeps, and writes HDF5 (observables) and XML (snapshots) outputs.

### 3.2.1 Build & Installation

The command-line interface requires HDF5  $\geq 1.14.3$  (older HDF5 versions may work, but were not tested). The core requires C++23, CMake  $\geq 3.23$ , and Boost  $\geq 1.87$  (older Boost versions may work, but were not tested). To compile it, run:

```

335 cmake -S . -B build -DCMAKE_BUILD_TYPE=Release \
336 -DPARATORIC_ENABLE_NATIVE_OPT=ON -DPARATORIC_LINK_MPI=OFF \
337 -DPARATORIC_BUILD_TESTS=ON -DPARATORIC_BUILD_CLI=ON
338 cmake --build build -jN
339 ctest --test-dir build -jN --output-on-failure
340 cmake --install build

```

Replace N with the number of cores to use, e.g. `-j4` for 4 cores.

- `-DCMAKE_BUILD_TYPE=Release`. Only set to Debug if you're a developer.
- `-DCMAKE_INSTALL_PREFIX`. By default, executables install to `${CMAKE_SOURCE_DIR}/${CMAKE_INSTALL_BINDIR}/`, headers to `${CMAKE_INSTALL_INCLUDEDIR}/paratoric`, and static libraries to `${CMAKE_SOURCE_DIR}/${CMAKE_INSTALL_LIBDIR}/`. The Python scripts expect `${CMAKE_SOURCE_DIR}/bin/`; this directory always contains the `paratoric` executable. To install into a custom directory, pass it via `-DCMAKE_INSTALL_PREFIX`, e.g. `-DCMAKE_INSTALL_PREFIX=/your/custom/directory/`.
- `-DPARATORIC_EXPORT_COMPILE_COMMANDS=ON`. Export `compile_commands.json` for tooling.
- `-DPARATORIC_LINK_MPI=OFF`. Link the core to MPI, required on some clusters. The core itself does not need MPI.
- `-DPARATORIC_ENABLE_NATIVE_OPT=OFF`. Turn on `-march=native` on GCC and Clang.

- 356 • -DPARATORIC\_ENABLE\_AVX2=OFF. Enable AVX2 (Haswell New Instructions). Requires  
357 a CPU which supports AVX2.
- 358 • -DPARATORIC\_BUILD\_TESTS=PROJECT\_IS\_TOP\_LEVEL. Compile the tests (recom-  
359 mended).
- 360 • -DPARATORIC\_BUILD\_CLI=PROJECT\_IS\_TOP\_LEVEL Required for both the C++ and  
361 Python command line interface.

## 362 Global options

Long flag	Short	Type	Description
--simulation	-sim	string	Simulation mode: <code>etc_sample</code> , <code>etc_hysteresis</code> , <code>etc_thermalization</code> .
--N_samples	-Ns	int	Number of recorded samples.
--N_thermalization	-Nth	int	Thermalization (warmup) steps.
--N_between_samples	-Nbs	int	Steps between samples (thinning).
--beta	-bet	double	Inverse temperature $\beta = 1/T$ .
--mu_constant	-muc	double	Star-term coupling $\mu$ .
--J_constant	-Jc	double	Plaquette coupling $J$ .
--h_constant	-hc	double	Field $h$ .
--lmbda_constant	-lmbdac	double	Field $\lambda$ .
--h_constant_therm	-hct	double	Thermalization value for $h$ (used if custom therm).
--lmbda_constant_therm	-lmbdact	double	Thermalization value for $\lambda$ .
--h_hysteresis	-hhys	list<double>	Hysteresis schedule for $h$ (space-separated). Length must match <code>lmbdahys</code> .
--lmbda_hysteresis	-lmbdahys	list<double>	Hysteresis schedule for $\lambda$ . Length must match <code>hhys</code> .
--N_resamples	-Nr	int	Bootstrap resamples (error bars).
--custom_therm	-cth	bool	Use thermalization values (0/1).
--observables	-obs	list<string>	Measured observables (space-separated).
--seed	-s	int	PRNG seed; 0 means random seed.
--basis	-bas	char	Spin basis: 'x' or 'z'.
--lattice_type	-lat	string	Lattice type (e.g. square, cubic, ...).
--system_size	-L	int	Linear lattice size (per dimension).
--boundaries	-bound	string	periodic or open.
--default_spin	-dsp	int	Initial link spin (+1 or -1).
--output_directory	-outdir	path	Output directory path.
--folder_name	-fn	string	Subfolder (of output directory) name for single run.
--folder_names	-fns	list<string>	Subfolders (of output directory) for hysteresis steps. Length must match <code>lmbdahys</code> .
--snapshots	-snap	bool	Save snapshots into specified subfolders of output directory.
--full_time_series	-fts	bool	Save full time series toggle.
--process_index	-procid	int	Process identifier (logging/debug).

## 364 3.2.2 etc\_sample

365 Runs a production measurement pass with the supplied configuration.

Listing 2: Example usage

366

```

367 ./paratoric -sim etc_sample -Ns 2000 -Nth 5000 -Nbs 10 -Nr 1000 -bet 16.0 -
368     muc 1 -Jc 1 -hc 0.2 -lmbdac 0.0 -obs energy plaquette_z anyon_count -bas
369     z -lat square -L 16 -bound periodic -dsp 1 -outdir ./runs/sample -snap
370     =0 -fcs=1
371

```

### 372 3.2.3 etc\_hysteresis

373 Runs a parameter sweep where the last state of step  $i$  initializes step  $i+1$ . Provide `--h_hy`  
374 `steresis` and `--lmbda_hysteresis` as space-separated lists, and `--folder_names` for  
375 per-step outputs.

Listing 3: Example usage

```

376 ./paratoric -sim etc_hysteresis -Ns 1000 -Nth 2000 -Nbs 50 -Nr 500 -bet
377     12.0 -muc 1 -Jc 1 -lmbdahys 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 -hhys 0.0 0.1
378     0.2 0.3 0.2 0.1 0.0 -obs energy fredenhagen_marcu -bas x -lat square -L
379     12 -bound periodic -dsp 1 -outdir ./runs/hys -fns step0 step1 step2
380     step3 step4 step5 step6
381
382

```

### 383 3.2.4 etc\_thermalization

384 Performs thermalization only (no production sampling).

Listing 4: Example usage

```

385 ./paratoric -sim etc_thermalization -Ns 0 -Nth 5000 -Nbs 10 -Nr 500 -bet
386     10.0 -muc 1 -Jc 1 -hc 0.3 -lmbdac 0.1 -hct 0.4 -lmbdact 0.2 -cth 1 -obs
387     energy anyon_density -bas z -lat square -L 10 -bound open -dsp 1 -outdir
388     ./runs/therm -snap=1
389
390

```

### 391 3.2.5 HDF5 structure

392 The output HDF5 has the structure `simulation/results/acc_ratio` for an array of the  
393 acceptance weights (only for thermalization), `simulation/results/observable_name/`  
394 `series` for the time series (if it was enabled), and `simulation/results/observable_`  
395 `name/{mean,mean_error,binder,binder_error,autocorrelation_time}`. For the  
396 regular sampling, `mean`, `mean_error`, `binder`, `binder_error` and `autocorrelation_`  
397 `time` contain doubles. For the hysteresis, they contain an array of values for the hysteresis  
398 parameters (in the order of `h_hysteresis` and `lmbda_hysteresis`).

## 399 3.3 C interface

400 The C interface enables users to use a ParaToric public header from within another C project  
401 or another programming language that supports a C-style interface. The C interface exposes  
402 a stable ABI. It mirrors the C++ interface. Include the public header `#include<paratoric/`  
403 `mcmc/extended_toric_code_c.h>`. All functions return `ptc_status_t`.

### 404 3.3.1 Build & Installation

405 The code can be compiled in exactly the same fashion as for the C++ interface.



406 **CMake usage (installed package).**

```
407 cmake_minimum_required(VERSION 3.23)
408 project(my_qmc_c C)
409 find_package(paratoric CONFIG REQUIRED) # provides paratoric::core
410 add_executable(cdemo main.c)
411 target_link_libraries(cdemo PRIVATE paratoric::core)
```

412 **CMake usage (as subdirectory).**

```
413 add_subdirectory(deps/paratoric)
414 add_executable(cdemo main.c)
415 target_link_libraries(cdemo PRIVATE paratoric::core)
```

### 416 3.3.2 Status & error handling

	Name	C type/values	Meaning
	ptc_status_t	{PTC_STATUS_OK=0, PTC_STATU S_INVALID_ARGUMENT=1, PTC_S TATUS_RUNTIME_ERROR=2, PTC_ STATUS_NO_MEMORY=3, PTC_STA TUS_INTERNAL_ERROR=4}	Return code of every API call.
417	ptc_last_error()	const char*	Thread-local error string. Valid until next call.

### 418 3.3.3 Opaque handle

419 Create and destroy the interface instance. Use ptc\_create(ptc\_handle\_t \*\*out) and p  
420 tc\_destroy(ptc\_handle\_t \*h).

### 421 3.3.4 Configuration type

422 Top-level ptc\_config\_t aggregates four nested specs. Field names mirror the C++ Config.

423 **Top-level configuration: ptc\_config\_t**

	Field	Type	Purpose
	sim	ptc_sim_spec_t	Monte Carlo parameters.
424	params	ptc_param_spec_t	Hamiltonian parameters.
	lat	ptc_lat_spec_t	Lattice parameters.
	out	ptc_out_spec_t	Output paths and snapshot toggle.

425 **Simulation specification (config.sim)**

Field	Type	Meaning
N_samples	int	Number of snapshots.
N_thermalization	int	Thermalization steps.
N_between_samples	int	Thinning between snapshots.
426 N_resamples	int	Bootstrap resamples.
custom_therm	bool	Custom thermalization schedule.
seed	int	PRNG seed (0 = random).
observables	const char* const*	Array of observable names (nullable).
N_observables	size_t	Length of observables.

427 **Parameter specification (config.params)**

Field	Type	Meaning
mu,h,J,lmbda	double	Couplings (star, electric, plaquette, gauge).
h_therm	double	Thermalization value for h if custom_therm=true.
lmbda_therm	double	Thermalization value for lmbda if custom_therm=true.
428 h_hys	const double*	Hysteresis schedule for h (nullable).
h_hys_len	size_t	Length of h_hys. Must match lmbda_hys_len.
lmbda_hys	const double*	Hysteresis schedule for lmbda (nullable).
lmbda_hys_len	size_t	Length of lmbda_hys. Must match h_hys_len.

429 **Lattice specification (config.lat)**

Field	Type	Meaning / Valid values
basis	char	Spin basis: 'x' or 'z'.
lattice_type	const char*	E.g. "triangular", "square", ...
430 system_size	int	Linear system size per dimension.
beta	double	Inverse temperature.
boundaries	const char*	"periodic" or "open".
default_spin	int	Initial link spin: +1 or -1.

431 **Output specification (config.out)**

Field	Type	Meaning
path_out	const char*	Single output directory (nullable).
paths_out	const char* const*	Output directories for hysteresis steps (nullable).
432 N_paths_out	size_t	Length of paths_out. Must match h_hys_len and lmbda_hys_len.
save_snapshots	bool	Toggle snapshot dumping.

433 **3.3.5 Return type**

434 All outputs are owned by the caller. Call `ptc_result_destroy(&r)` to free and zero.

Field	C type	Meaning
series	ptc_series_t	Time series (real/complex) of all requested observables, thermalization is excluded (except for thermalization simulation). Outer index = observable, inner index = time point.
acc_ratio	ptc_dvec_t	MC acceptance ratios for each update (thermalization).
mean, mean_std	ptc_dvec_t	Bootstrap mean and standard error (order as in observables).
binder, binder_std	ptc_dvec_t	Binder ratios and standard error (order as in observables).
435 tau_int	ptc_dvec_t	Integrated autocorrelation time (order as in observables).
series_hys	ptc_series_block_s_t	Hysteresis time series. Outer index = hysteresis parameters (order as in h_hys, lmbda_hys), middle index = observables (order as in observables), inner vector = time series.
mean_hys, mean_std_hys, binder_hys, binder_std_hys, tau_int_hys	ptc_dmat_t	Outer index = hysteresis parameters (order as in h_hys, lmbda_hys), inner index = observables (order as in observables).

### 436 3.3.6 Procedures (mirror the C++ API)

437 All fill a ptc\_result\_t \*out on success. Return PTC\_STATUS\_OK on success.

438 ptc\_get\_thermalization(ptc\_handle\_t \*h, const ptc\_config\_t \*cfg, ptc\_result\_t \*out) Run thermalization only. **Required fields:** cfg->lat.{basis,lattice\_type,system\_size,beta,boundaries,default\_spin}, cfg->params.{mu,h,J,lmbda}, cfg->sim.{N\_thermalization,N\_resamples,observables,N\_observables,seed}, cfg->out.{path\_out,save\_snapshots}.

443 ptc\_get\_sample(ptc\_handle\_t \*h, const ptc\_config\_t \*cfg, ptc\_result\_t \*out) Run a production measurement pass. **Required fields:** cfg->lat.{basis,lattice\_type,system\_size,beta,boundaries,default\_spin}, cfg->params.{mu,h,J,lmbda,h\_therm,lmbda\_therm}, cfg->sim.{N\_samples,N\_thermalization,N\_between\_samples,N\_resamples,custom\_therm,observables,N\_observables,seed}, cfg->out.{path\_out,save\_snapshots}.

449 ptc\_get\_hysteresis(ptc\_handle\_t \*h, const ptc\_config\_t \*cfg, ptc\_result\_t \*out) Run a hysteresis sweep over h\_hys and/or lmbda\_hys. The last state of step i initializes step i+1. **Required fields:** cfg->lat.{basis,lattice\_type,system\_size,beta,boundaries,default\_spin}, cfg->params.{mu,h\_hys,h\_hys\_len,J,lmbda\_hys,lmbda\_hys\_len}, cfg->sim.{N\_samples,N\_thermalization,N\_between\_samples,N\_resamples,observables,N\_observables,seed}, cfg->out.{paths\_out,N\_paths\_out,save\_snapshots}.

### 456 3.3.7 C usage example

Listing 5: C API - Minimal call

```

457 #include <stdio.h>
458 #include <math.h>
459 #include <paratoric/mcmc/extended_toric_code_c.h>
460
461
462 int main(void) {
463     ptc_handle_t* h = NULL;
464     if (ptc_create(&h) != PTC_STATUS_OK) { puts("create_failed"); return 1; }
465
466     ptc_lat_spec_t lat = {
467         .basis = 'z',
468         .lattice_type = "square",
469         .system_size = 16,
470         .beta = 8.0,
471         .boundaries = "periodic",
472         .default_spin = 1
473     };
474
475     ptc_param_spec_t ps = {
476         .mu = 1.0, .h = 0.2, .J = 1.0, .lmbda = 0.0,
477         .h_therm = NAN, .lmbda_therm = NAN,
478         .h_hys = NULL, .h_hys_len = 0,
479         .lmbda_hys = NULL, .lmbda_hys_len = 0
480     };
481
482     const char* obs[] = {"energy", "plaquette_z", "anyon_count"};
483     ptc_sim_spec_t sim = {
484         .N_samples = 0, /* thermalization-only initially */
485         .N_thermalization = 5000,
486         .N_between_samples = 10,
487         .N_resamples = 1000,
488         .custom_therm = false,
489         .seed = 12345,
490         .observables = obs,
491         .N_observables = sizeof(obs)/sizeof(obs[0])
492     };
493
494     ptc_out_spec_t outspec = {
495         .path_out = "runs/sample",
496         .paths_out = NULL, .N_paths_out = 0,
497         .save_snapshots = false
498     };
499
500     ptc_config_t cfg = { .sim = sim, .params = ps, .lat = lat, .out = outspec
501         };
502
503     ptc_result_t warm = {0};
504     ptc_status_t st = ptc_get_thermalization(h, &cfg, &warm);
505     if (st != PTC_STATUS_OK) { puts(ptc_last_error()); ptc_destroy(h); return
506         2; }
507     ptc_result_destroy(&warm);
508
509     cfg.sim.N_samples = 2000;
510     ptc_result_t res = {0};
511     st = ptc_get_sample(h, &cfg, &res);
512     if (st != PTC_STATUS_OK) { puts(ptc_last_error()); ptc_destroy(h); return
513         3; }

```

```

514
515 // use res.mean, res.tau_int, ...
516 ptc_result_destroy(&res);
517 ptc_destroy(h);
518 return 0;
519 }
520

```

521 **Memory rules.** You own all buffers in `ptc_result_t`. Call `ptc_result_destroy` once  
 522 per successful call.

### 523 3.4 Python bindings

524 ParaToric exposes a compiled Python extension module `_paratoric` with a submodule `ext`  
 525 `ended_toric_code`. The bindings convert C++ vectors into NumPy arrays and release the  
 526 global interpreter lock (GIL) while running the C++ kernels.

#### 527 3.4.1 Build & Installation

528 The core requires C++23, CMake  $\geq 3.23$ , and Boost  $\geq 1.87$  (older Boost versions may work,  
 529 but were not tested). The Python bindings require a Python installation with Numpy and  
 530 PyBind11 (tested with version 3.0.1). Pybind11 is included as a git submodule (you need to  
 531 pull it!). To compile the Python bindings, run:

```

532 cmake -S . -B build -DCMAKE_BUILD_TYPE=Release \
533 -DPARATORIC_ENABLE_NATIVE_OPT=ON -DPARATORIC_LINK_MPI=OFF \
534 -DPARATORIC_BUILD_TESTS=ON -DPARATORIC_BUILD_PYBIND=ON \
535 -DPython3_EXECUTABLE="$(which python)" -DPYBIND11_FINDPYTHON=ON \
536 -DPARATORIC_INSTALL_TO_SITE=ON -DPARATORIC_PIP_EDITABLE_INSTALL=ON
537 cmake --build build -jN
538 ctest --test-dir build -jN --output-on-failure
539 cmake --install build

```

540 Replace N with the number of cores to use, e.g. `-j4` for 4 cores.

- 541 • `-DCMAKE_BUILD_TYPE=Release`. Only set to Debug if you're a developer.
- 542 • `-DCMAKE_INSTALL_PREFIX`. By default, executables install to `${CMAKE_SOURCE_DIR}/${CMAKE_INSTALL_BINDIR}/`, headers to `${CMAKE_INSTALL_INCLUDEDIR}/`  
 543 `paratoric`, and static libraries to `${CMAKE_SOURCE_DIR}/${CMAKE_INSTALL_LIBDIR}/`. The Python scripts expect `${CMAKE_SOURCE_DIR}/bin/`; this directory always  
 544 contains the `paratoric` executable. To install into a custom directory, pass it via `-DCMAKE_INSTALL_PREFIX`, e.g. `-DCMAKE_INSTALL_PREFIX=/your/custom/directo`  
 545 `ry/`.
- 546 • `-DPARATORIC_EXPORT_COMPILE_COMMANDS=ON`. Export `compile_commands.json`  
 547 for tooling.
- 548 • `-DPARATORIC_LINK_MPI=OFF`. Link the core to MPI, required on some clusters. The  
 549 core itself does not need MPI.
- 550 • `-DPARATORIC_ENABLE_NATIVE_OPT=OFF`. Turn on `-march=native` on GCC and  
 551 Clang.
- 552 • `-DPARATORIC_ENABLE_NATIVE_OPT=OFF`. Turn on `-march=native` on GCC and  
 553 Clang.
- 554

- 555 • `-DPARATORIC_ENABLE_AVX2=OFF`. Enable AVX2 (Haswell New Instructions). Requires  
556 a CPU which supports AVX2.
- 557 • `-DPARATORIC_BUILD_TESTS=PROJECT_IS_TOP_LEVEL`. Compile the tests (recom-  
558 mended).
- 559 • `-DPARATORIC_BUILD_PYBIND=OFF`. Compile Python bindings.
- 560 • `-DDPARATORIC_INSTALL_TO_SITE=OFF`. Install ParaToric Python module to site pack-  
561 ages.
- 562 • `-DPARATORIC_PIP_EDITABLE_INSTALL=OFF`. Install ParaToric Python module via  
563 pip as an editable module.
- 564 • `-DPARATORIC_PIP_OFFLINE_INSTALL=OFF`. Turn on when installing to pip without  
565 internet access. Requires NumPy and setuptools.

### 566 3.4.2 Module layout

- 567 • `paratoric._paratoric`: compiled extension (PyBind11). Submodule: `extended_t`  
568 `oric_code`.
- 569 • `paratoric.extended_toric_code`: convenient alias
- 570 • Running `python -m paratoric` enters the package entry point (`__main__.py`).

### 571 NumPy return formats

572 All time series with potentially complex values are returned as `complex128`. Real observables  
573 appear with zero imaginary part. Shapes are documented in the function references below.

### 574 API reference (`paratoric.extended_toric_code`)

575 `get_thermalization(...)` Run only the warmup and return per-snapshot observables  
576 and MC acceptance ratios. Internally converts `std::variant<complex<double>, double`  
577 `>` to `complex128` and `std::vector<double>` to `float64` arrays. The GIL is released while  
578 the C++ routine executes.

	Parameter	Type / default	Meaning
	<code>N_thermalization</code>	<code>int</code>	Warmup steps.
	<code>N_resamples</code>	<code>int=1000</code>	Bootstrap resamples.
	<code>observables</code>	<code>list[str]</code>	Names per snapshot.
	<code>seed</code>	<code>int=0</code>	PRNG seed (0 $\Rightarrow$ random).
	<code>mu, h, J, lambda</code>	<code>float</code>	Hamiltonian parameters.
	<code>basis</code>	<code>{'x', 'z'}='x'</code>	Spin eigenbasis.
579	<code>lattice_type</code>	<code>str</code>	E.g. "triangular", "square", ...
	<code>system_size</code>	<code>int</code>	Linear size per dimension.
	<code>beta</code>	<code>float</code>	Inverse temperature.
	<code>boundaries</code>	<code>str="periodic"</code>	Boundary condition.
	<code>default_spin</code>	<code>int=1</code>	Initial link spin (+1/-1).
	<code>save_snapshots</code>	<code>bool=false</code>	Enable snapshot files.
	<code>path_out</code>	<code>path None=None</code>	Output directory (if saving).

580 **Returns:** (`series, acc_ratio`) `series`: `ndarray(complex128)` of shape (`n_ob`  
581 `s, N_thermalization`); `acc_ratio`: `ndarray(float64)` of shape (`N_thermalizati`  
582 `on,`).

583 `get_sample(...)` Run thermalization and production sampling; return series and boot-  
 584 strap statistics. Converts nested C++ containers to NumPy arrays and releases the GIL during  
 585 computation.

Parameter	Type / default	Meaning
<code>N_samples</code>	<code>int</code>	Stored samples per observable.
<code>N_thermalization</code>	<code>int</code>	Warmup steps before sampling.
<code>N_between_samples</code>	<code>int</code>	Thinning between samples.
<code>N_resamples</code>	<code>int=1000</code>	Bootstrap resamples.
<code>custom_therm</code>	<code>bool=false</code>	Use <code>h_therm</code> , <code>lmbda_therm</code> during warmup.
<code>observables</code>	<code>list[str]</code>	Names per snapshot.
<code>seed</code>	<code>int=0</code>	PRNG seed (0 $\Rightarrow$ random).
586 <code>mu, h, J, lmbda</code>	<code>float</code>	Hamiltonian parameters.
<code>h_therm, lmbda_therm</code>	<code>float=0</code>	Warmup parameters if custom therm.
<code>basis</code>	<code>{'x', 'z'}='x'</code>	Spin eigenbasis.
<code>lattice_type, system_size, beta</code>	<code>str, int, float</code>	Lattice and temperature.
<code>a</code>		
<code>boundaries, default_spin</code>	<code>str, int=("periodic", 1)</code>	BC and initial spin.
<code>save_snapshots, path_out</code>	<code>bool=False, path   None=None</code>	Optional I/O.

587 **Returns:** tuple of six arrays `series (complex128): (n_obs, N_samples); mean,`  
 588 `mean_std, binder, binder_std, tau_int (float64): each (n_obs,).`

589 `get_hysteresis(...)` Run a sweep where each step uses the previous state as its ini-  
 590 tial condition. Returns stacked arrays across steps; path handling validates per-step output  
 591 directories when saving snapshots.

Parameter	Type / default	Meaning
<code>N_samples, N_thermalization, N_between_samples</code>	<code>int, int, int</code>	Cadence per step.
<code>N_resamples</code>	<code>int=1000</code>	Bootstrap resamples.
<code>observables</code>	<code>list[str]</code>	Names per snapshot.
<code>seed</code>	<code>int=0</code>	PRNG seed.
<code>mu, J</code>	<code>float</code>	Star and plaquette couplings.
592 <code>h_hys, lmbda_hys</code>	<code>list[float]</code>	Hysteresis values, length must match.
<code>basis</code>	<code>{'x', 'z'}='x'</code>	Spin basis.
<code>lattice_type, system_size, beta</code>	<code>str, int, float</code>	Lattice and temperature.
<code>boundaries, default_spin</code>	<code>str, int=("periodic", 1)</code>	BC and initial spin.
<code>save_snapshots</code>	<code>bool=false</code>	Enable stepwise I/O.
<code>paths_out</code>	<code>list[path]   None=None</code>	Output path per step (size must match <code>h_hys</code> if saving).

593 **Returns:** tuple of six arrays `series3d (complex128): (n_steps, n_obs, N_sam`  
 594 `ples); mean2d, std2d, binder2d, binder_std2d, tau2d (float64): each (n_steps,`  
 595 `n_obs).` The number of steps equals `len(h_hys)` (and `len(lmbda_hys)`).

### 596 Array dtypes and shapes (summary)

Function	Name / dtype	Shape
get_thermalization	series (complex128) acc_ratio (float64)	(n_obs, N_thermalization) (N_thermalization,)
get_sample	series (complex128) mean, mean_std, binder, binder_std, tau_int (float64)	(n_obs, N_samples) each (n_obs,)
get_hysteresis	series3d (complex128) mean2d, std2d, binder2d, binder_std2d, tau2d (float64)	(n_steps, n_obs, N_samples) each (n_steps, n_obs)

### 598 Notes on performance

599 The bindings release the global interpreter lock (GIL) during heavy compute (py::gil\_scoped\_release), enabling multi-threaded C++ execution if the backend uses threads or when  
600 calling from multiprocessing workers. Conversions handle 1D/2D/3D containers and enforce  
601 consistent inner lengths before copying to NumPy.  
602

### 603 3.4.3 Usage example

Listing 6: Importing and calling from Python

```

604 >>> import numpy as np
605 >>> from paratoric import extended_toric_code as etc
606 >>> series, acc_ratio = etc.get_thermalization(
607 ... N_thermalization=2000, N_resamples=500,
608 ... observables=["energy", "plaquette_z", "anyon_count"],
609 ... seed=0, mu=1.0, h=0.2, J=1.0, lmbda=0.0,
610 ... basis='z', lattice_type="square", system_size=16, beta=8.0,
611 ... boundaries="periodic", default_spin=1,
612 ... save_snapshots=False, path_out=None)
613 >>> series.shape, series.dtype
614 ((3, 2000), dtype('complex128'))
615 >>> out = etc.get_sample(
616 ... N_samples=1000, N_thermalization=5000, N_between_samples=10,
617 ... N_resamples=1000, custom_therm=False,
618 ... observables=["energy", "plaquette_z"],
619 ... seed=0, mu=1.0, h=0.2, h_therm=0.0,
620 ... J=1.0, lmbda=0.0, lmbda_therm=0.0,
621 ... basis='z', lattice_type="square", system_size=16, beta=8.0,
622 ... boundaries="periodic", default_spin=1,
623 ... save_snapshots=False, path_out=None)
624 >>> (series_s, mean, mean_std, binder, binder_std, tau_int) = out
625
626

```

### 627 3.5 Python command-line interface

628 ParaToric ships a Python command-line interface /python/cli/paratoric.py that orches-  
629 trates C++ backends, runs sweeps, and writes HDF5/XML outputs. It requires [NumPy](#), [mat-](#)  
630 [plotlib](#), and [h5py](#).



### 3.5.1 Build & Installation

The command-line interface requires HDF5  $\geq 1.14.3$  (older HDF5 versions may work, but were not tested). The core requires C++23, CMake  $\geq 3.23$ , and Boost  $\geq 1.87$  (older Boost versions may work, but were not tested). To compile it, run:

```
cmake -S . -B build -DCMAKE_BUILD_TYPE=Release \
-DPARATORIC_ENABLE_NATIVE_OPT=ON -DPARATORIC_LINK_MPI=OFF \
-DPARATORIC_BUILD_TESTS=ON -DPARATORIC_BUILD_CLI=ON
cmake --build build -jN
ctest --test-dir build -jN --output-on-failure
cmake --install build
```

Replace N with the number of cores to use, e.g. -j4 for 4 cores.

- -DCMAKE\_BUILD\_TYPE=Release. Only set to Debug if you're a developer.
- -DCMAKE\_INSTALL\_PREFIX. By default, executables install to  $\{\text{CMAKE\_SOURCE\_DIR}\}/\{\text{CMAKE\_INSTALL\_BINDIR}\}/$ , headers to  $\{\text{CMAKE\_INSTALL\_INCLUDEDIR}\}/\text{paratoric}$ , and static libraries to  $\{\text{CMAKE\_SOURCE\_DIR}\}/\{\text{CMAKE\_INSTALL\_LIBDIR}\}/$ . The Python scripts expect  $\{\text{CMAKE\_SOURCE\_DIR}\}/\text{bin}/$ ; this directory always contains the paratoric executable. To install into a custom directory, pass it via -DCMAKE\_INSTALL\_PREFIX, e.g. -DCMAKE\_INSTALL\_PREFIX=/your/custom/directory/.
- -DPARATORIC\_EXPORT\_COMPILE\_COMMANDS=ON. Export compile\_commands.json for tooling.
- -DPARATORIC\_LINK\_MPI=OFF. Link the core to MPI, required on some clusters. The core itself does not need MPI.
- -DPARATORIC\_ENABLE\_NATIVE\_OPT=OFF. Turn on -march=native on GCC and Clang.
- -DPARATORIC\_ENABLE\_AVX2=OFF. Enable AVX2 (Haswell New Instructions). Requires a CPU which supports AVX2.
- -DPARATORIC\_BUILD\_TESTS=PROJECT\_IS\_TOP\_LEVEL. Compile the tests (recommended).
- -DPARATORIC\_BUILD\_CLI=PROJECT\_IS\_TOP\_LEVEL Required for both the C++ and Python command line interface.

662 **General options**

Long flag	Short	Description
--help	-h	Show help and exit.
--simulation	-sim	Simulation type selector.
--N_thermalization	-Nth	Thermalization steps (proposed updates).
--N_samples	-Ns	Number of samples/snapshots.
--N_between_steps	-Nbs	Steps between successive samples (thinning).
--N_resamples	-Nr	Bootstrap resamples.
--custom_therm	-cth	Use thermalization values for $h, \lambda$ (0 or 1).
--observables	-obs	Space-separated list, e.g. fredenhagen_marcu percolation_probability energy.
--seed	-seed	PRNG seed; 0 means random seed.
663 --mu_constant	-muc	Value of $\mu$ .
--J_constant	-Jc	Value of $J$ .
--h_constant	-hc	Value of $h$ .
--h_constant_therm	-hct	Thermalization value of $h$ .
--lambda_constant	-lmbdac	Value of $\lambda$ .
--lambda_constant_therm	-lmbdact	Thermalization value of $\lambda$ .
--output_directory	-outdir	Output directory.
--snapshots	-snap	Save snapshots toggle (0/1).
--full_time_series	-fts	Save full time series toggle (0/1).
--processes	-proc	Logical CPU count for Python multiprocessing. 0 means all available cores. Negative numbers $-x$ mean use all cores minus $x$ . Default is -4.

664 **Lattice-specific options**

Long flag	Short	Description
--help	-h	Show help and exit.
--basis	-bas	Spin basis: x or z.
--lattice_type	-lat	square, cubic, triangular, honeycomb, ...
665 --system_size	-L	Linear size; in 2D, 30 yields a $30 \times 30$ lattice (unit cells).
--temperature	-T	Temperature $T = 1/\beta > 0$ .
--boundaries	-bound	periodic or open.
--default_spin	-dsp	Initial edge spin: 1 or -1.

666 The command line interface offers several sweep modes. All are embarrassingly parallel;  
 667 set `--processes` close to the number of steps when possible.

668 **3.5.2  $T$ -sweep**

669 Runs `T_steps` independent Markov chains for evenly spaced temperatures in  $[T_{\text{lower}}, T_{\text{upper}}]$   
 670 and plots all requested observables.

Listing 7: Example usage

671

```

672 python3 ./python/cli/paratoric.py -sim etc_T_sweep -Ns 1000 -muc 1 -Nth
673     2000 -Nbs 100 -Tl 0.5 -Tu 5 -Ts 30 -hc 0.1 -Jc 1 -lmbdac 0.1 -Nr 1000 -
674     obs percolation_strength percolation_probability
675     plaquette_percolation_probability largest_cluster string_number energy
676     energy_h energy_mu energy_J energy_lambda sigma_x sigma_z star_x
677     plaquette_z staggered_imaginary_times delta anyon_count anyon_density
678     fredenhagen_marcu sigma_x_susceptibility sigma_z_susceptibility -s 0 -
679     bas x -lat square -L 4 -bound periodic -dsp 1 -outdir /path/to/out

```

### 681 Sweep-specific flags.

Long flag	Short	Description
--simulation	-sim	Use etc_T_sweep.
682 --T_lower	-Tl	Lower bound of $T$ .
--T_upper	-Tu	Upper bound of $T$ .
--T_steps	-Ts	Number of temperatures between bounds.

### 683 3.5.3 h-sweep

684 Runs  $h\_steps$  independent chains in parallel for evenly spaced  $h$  in  $[h\_lower, h\_upper]$ .

Listing 8: Example usage

```

685 python3 ./python/cli/paratoric.py -sim etc_h_sweep -Ns 1000 -muc 1 -Nth
686     2000 -Nbs 100 -hl 0.1 -hu 0.5 -hs 8 -T 0.03 -Jc 1 -lmbdac 0.2 -Nr 1000 -
687     obs percolation_strength percolation_probability
688     plaquette_percolation_probability largest_cluster string_number energy
689     energy_h energy_mu energy_J energy_lambda sigma_x sigma_z star_x
690     plaquette_z staggered_imaginary_times delta anyon_count anyon_density
691     fredenhagen_marcu sigma_x_susceptibility sigma_z_susceptibility -s 0 -
692     bas x -lat square -L 6 -bound periodic -dsp 1 -outdir /path/to/out
693

```

### 695 Sweep-specific flags.

Long flag	Short	Description
--simulation	-sim	Use etc_h_sweep.
696 --h_lower	-hl	Lower bound of $h$ .
--h_upper	-hu	Upper bound of $h$ .
--h_steps	-hs	Number of field steps between bounds.

### 697 3.5.4 $\lambda$ -sweep

698 Runs  $lambda\_steps$  independent chains in parallel for evenly spaced  $\lambda$  in  $[\lambda\_lower, \lambda\_upper]$ .

Listing 9: Example usage

```

699 python3 ./python/cli/paratoric.py -sim etc_lambda_sweep -Ns 1000 -muc 1 -Nth
700     2000 -Nbs 100 -lmbdal 0.01 -lmbdau 1.0 -lmbdas 15 -T 0.1 -hc 0.3 -Jc 1
701     -Nr 1000 -obs percolation_strength percolation_probability
702     plaquette_percolation_probability largest_cluster string_number energy
703     energy_h energy_mu energy_J energy_lambda sigma_x sigma_z star_x
704     plaquette_z staggered_imaginary_times delta anyon_count anyon_density
705     fredenhagen_marcu sigma_x_susceptibility sigma_z_susceptibility -s 0 -
706     bas x -lat square -L 4 -bound periodic -dsp 1 -outdir /path/to/out
707

```

## 709 Sweep-specific flags.

	Long flag	Short	Description
	--simulation	-sim	Use etc_lambda_sweep.
710	--lambda_lower	-lmbdal	Lower bound of $\lambda$ .
	--lambda_upper	-lmbdau	Upper bound of $\lambda$ .
	--lambda_steps	-lmbdas	Number of field steps between bounds.

## 711 3.5.5 o-sweep

712 Runs Theta\_steps independent chains in parallel along a circle in  $(\lambda, h)$  centered at  $(\text{lambda\_constant}, \text{h\_constant})$  with radius radius, for angles  $\Theta \in [\Theta_{\text{lower}}, \Theta_{\text{upper}}]$  (angles measured anti-clockwise from the  $\lambda$ -axis).

Listing 10: Example usage

```
715 python3 ./python/cli/paratoric.py -sim etc_circle_sweep -Ns 1000 -muc 1 -
716   Nth 2000 -Nbs 100 -lmbdac 0.4 -rad 0.3 -Thl 0 -Thu 3.141 -Ths 15 -T 0.1
717   -hc 0.4 -Jc 1 -Nr 1000 -obs percolation_strength percolation_probability
718   plaquette_percolation_probability largest_cluster string_number energy
719   energy_h energy_mu energy_J energy_lambda sigma_x sigma_z star_x
720   plaquette_z staggered_imaginary_times delta anyon_count anyon_density
721   fredenhagen_marcu sigma_x_susceptibility sigma_z_susceptibility -s 0 -
722   bas x -lat square -L 4 -bound periodic -dsp 1 -outdir /path/to/out
723
```

## 725 Sweep-specific flags.

	Long flag	Short	Description
	--simulation	-sim	Use etc_circle_sweep.
	--lambda_constant	-lmbdac	Circle center in $\lambda$ .
	--h_constant	-hc	Circle center in $h$ .
726	--radius	-rad	Circle radius.
	--Theta_lower	-Thl	Lower bound of $\Theta$ .
	--Theta_upper	-Thu	Upper bound of $\Theta$ .
	--Theta_steps	-Ths	Number of angles between bounds.

## 727 3.5.6 Hysteresis-sweep

728 Use the hysteresis schedule specified in hhys and lmbdahys. This mode will run two Markov  
729 chains, one in the original parameter order specified in hhys and lmbdahys, and one with a  
730 reversed parameter order, i.e., it calculates both branches of the hysteresis loop.

Listing 11: Example usage

```
731 python3 ./python/cli/paratoric.py -sim etc_hysteresis -Nbs 5000 -Ns 10000 -
732   muc 1 -Nth 20000 -hhys 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 -T
733   0.1 -Jc -1 -lmbdahys 0.5 0.525 0.55 0.575 0.6 0.625 0.65 0.675 0.7 0.725
734   0.75 -Nr 1000 -obs plaquette_percolation_probability
735   percolation_strength percolation_probability largest_cluster
736   string_number energy energy_h energy_mu energy_J energy_lambda sigma_x
737   sigma_z star_x plaquette_z staggered_imaginary_times delta anyon_count
738   anyon_density fredenhagen_marcu sigma_x_susceptibility
739   sigma_z_susceptibility -s 0 -bas z -lat square -L 4 -bound periodic -dsp
740   1 -outdir /path/to/out
741
```

### 743 Sweep-specific flags.

Long flag	Short	Description
--simulation	-sim	Use etc_hysteresis.
--lmbda_hysteresis	-lmbdahys	Hysteresis schedule for $\lambda$ . Length must match hhys.
--h_hysteresis	-hhys	Hysteresis schedule for $h$ . Length must match lmbdahys.

### 745 3.5.7 Thermalization

746 Runs repetitions independent chains in parallel and reports observables and MC accep-  
 747 tance ratios every step, averaged over chains.

Listing 12: Example usage

```

748 python3 ./python/cli/paratoric.py -sim etc_thermalization -muc 1 -Nth 2000
749 -reps 10 -lmbdac 2 -T 0.1 -hc 0.3 -Jc 1 -Nr 1000 -obs
750 percolation_strength percolation_probability
751 plaquette_percolation_probability largest_cluster string_number energy
752 energy_h energy_mu energy_J energy_lambda sigma_x sigma_z star_x
753 plaquette_z staggered_imaginary_times delta anyon_count anyon_density
754 fredenhagen_marcu sigma_x_susceptibility sigma_z_susceptibility -s 0 -
755 bas x -lat square -L 4 -bound periodic -dsp 1 -outdir /path/to/out
756

```

### 758 Sweep-specific flags.

Long flag	Short	Description
--simulation	-sim	Use etc_thermalization.
--repetitions	-reps	Number of Markov chains to average.

## 760 4 Using ParaToric

### 761 4.1 Monte Carlo Updates

762 There is no need for the user to explicitly call specific updates or interact with internal C++  
 763 classes when using the documented interfaces. Internally, we use all five updates described  
 764 in the original algorithm by Wu, Deng, and Prokof'ev [5]. These must furthermore be sup-  
 765 plemented with the following two updates: Because for high temperatures and for zero off-  
 766 diagonal fields the spin at imaginary time  $0 = \beta$  cannot be flipped, we allow for flipping the  
 767 spin on the entire imaginary axis on one bond or on a plaquette (star) in the  $\hat{\sigma}^x$ -basis ( $\hat{\sigma}^z$ -  
 768 basis).

769 These updates only change the energy terms diagonal in the given basis and is trivial when  
 770 caching the total integrated diagonal energy (the update locally flips the sign of the total in-  
 771 tegrated potential energy). Another advantage is that integrated autocorrelation times for  
 772 observables diagonal in the given basis improve even in regimes that were previously accessi-  
 773 ble. All seven updates are equally likely to be proposed, and we use a 64-bit Mersenne-Twister  
 774 for pseudorandom numbers [12] with the ability to externally set the seed. Some updates  
 775 have early exits for input parameters for which they will always be rejected.

## 4.2 Monte Carlo Diagnostics

There are two compilation modes, Release and Debug. In production runs, one should always use the Release mode; however, it still gives the user enough information to diagnose sampling problems without severe performance impacts.

### 4.2.1 Thermalization mode

We provide thermalization routines which should be used before production runs to ensure proper thermalization (also known as burn-in). Thermalization times can vary drastically between different observables and initial conditions. We provide an example of sufficient and insufficient thermalization in Fig. 2. We recommend using the provided Python command-line interface, which will also plot the thermalization of all measured observables for the user.

In thermalization runs, we also return the Monte Carlo acceptance ratio of every update. This can also be used to diagnose freezing (in the measurement phase, use the integrated autocorrelation time instead), e.g., when the acceptance ratio is always identical and/or very low.

In case one suspects experiencing a serious sampling problem, we recommend recompiling the project in the Debug mode, which provides a wide array of runtime debug information about the proposed steps, acceptance ratios, and intermediate results. However, do not use the Debug mode in production runs, as it negatively impacts performance.

### 4.2.2 Integrated autocorrelation time

When measuring observables, we first thermalize the system with `N_thermalization` steps, then measure `N_samples` times with `N_between_samples` steps between measurements. The normalized autocorrelation function  $\rho_O(k)$  of an observable  $O_k$  (observable  $O$  measured at time  $k$ ) applied to a discrete time series of length  $N$  is given by:

$$\rho_O(k) = \frac{C(k)}{C(0)}, \quad C(k) = \frac{1}{N-k} \sum_{i=0}^{N-k-1} (O_i - \bar{O})(O_{i+k} - \bar{O}), \quad \bar{O} = \frac{1}{N} \sum_{i=0}^{N-1} O_i. \quad (7)$$

It is a statistical measure of the correlations between measurements of observable  $O$  at times  $i$  and  $i+k$ .<sup>2</sup> We define the *integrated autocorrelation time*

$$\tau_{\text{int}}^O = \frac{1}{2} + \sum_{k \geq 1} \rho_O(k). \quad (8)$$

Large  $\tau_{\text{int}}$  are generally undesirable since they increase error bars and can lead to bias. In case of perfect sampling, we would have  $\rho_O(0) = 1$  and  $\rho_O(k) = 0 \forall k \geq 1$ , i.e., each measurement is only correlated with itself but not with other measurements and  $\tau_{\text{int}} = 1/2$ . In practice, this is usually not feasible, and we have to work with a finite autocorrelation time  $\tau_{\text{int}} > 1/2$ .

When using ParaToric, we strongly recommend monitoring  $\tau_{\text{int}}$  for all simulations and all observables. It is automatically calculated for every observable based on the full time series. As a rule of thumb, the autocorrelation is fine as long as  $\tau_{\text{int}} \ll N_{\text{samples}}$ , otherwise it leads to bias and seriously underestimated error bars. In the vicinity of phase transitions,  $\tau_{\text{int}}$  dramatically increases (“critical slowing down”) [13]. Importantly,  $\tau_{\text{int}}$  can differ vastly between different observables! If the autocorrelation is too high, increase the number of steps between samples. In more complicated cases, one may need to adapt the update proposal distributions and/or the updates themselves as a last resort.

<sup>2</sup>In ParaToric, the autocorrelation function is calculated efficiently using fast Fourier transforms.

It is also important to mention that ParaToric only computes a statistical *estimate* of  $\tau_{\text{int}}$ . Many factors determine how accurate this estimate is, and crucially, the system needs to be properly *thermalized*. In principle, one can use  $\tau_{\text{int}}$  in the way it is computed above directly for calculating error bars of correlated time series; however, ParaToric uses a more robust bootstrapping approach.

### 4.2.3 Error bars

ParaToric applies the stationary bootstrap [14–16] for all error bars, thus capturing autocorrelation effects. Large  $\tau_{\text{int}}$  will lead to worse error bars. The only parameter that the user can change is the number of bootstrap resamples `N_between_steps`. The default is 1000, which is enough in most cases. Note that a too low value of `N_between_steps` increases the relative computational cost of performing the measurements, which may negatively affect the code efficiency at no statistical gain. If the error bars are too large, either the number of samples is too low (in which case one should increase `N_samples`) or the autocorrelation is too large (in which case one could additionally increase `N_between_steps`).

## 4.3 Tips & tricks

### 4.3.1 Probing ground state physics

The algorithm implemented by ParaToric fundamentally requires a finite temperature  $T > 0$ . However, in QMC simulations, there is always a finite-size energy gap (the difference between the energy of the ground state and the first excited state). Additionally, some phases like the topological ground state of the toric code have a physical bulk gap (even at  $L \rightarrow \infty$ ). As long as the temperature is well below the total gap, we are exponentially close to the ground state. Usually, a temperature  $T \sim 1/L$  suffices for the toric code although other situations may arise.

### 4.3.2 Probing first-order transitions

ParaToric provides functionalities to probe weak and strong first-order phase transitions. The hysteresis mode can be used to probe hysteresis loops in the vicinity of strong first-order phase transitions, by repeating the simulation two times and mirroring the order of the parameters in `h_hysteresis` and `lmbda_hysteresis`. Weak first-order transitions can be detected by plotting a time series histogram of an observable (it exhibits a double-peak structure). Both approaches have been used in the context of the toric code [11].

### 4.3.3 Choosing the basis

Sometimes, one can work in both the  $\hat{\sigma}^x$  and the  $\hat{\sigma}^z$ -basis. The performance can vary drastically! Generally, the  $\hat{\sigma}^x$ -basis is more efficient for  $h/J > \lambda/\mu$  and vice versa.

### 4.3.4 Choosing `N_thermalization`

Based on our experience, we can say that  $N_{\text{thermalization}} = 500L^d/T$  is a sensible choice for small fields, where  $d$  is the dimensionality of the system. Nevertheless, one should make use of the provided tools to benchmark thermalization, see Sec. 4.2, and rather err on the side of safety.

### 4.3.5 Choosing `N_samples`

Neglecting autocorrelation effects, the error of an observable  $\Delta O$  scales as  $\Delta O \sim 1/\sqrt{N_{\text{samples}}}$ . More samples are, in principle, always better and lead to lower error



bars. Smoothness of a curve of statistical results also requires that error bars be small in relation to the parameter grid size. If one increases the parameter resolution (e.g., in the field  $h$ ), then one typically also increases `N_samples`.

#### 4.3.6 Choosing `N_between_samples`

The optimal choice for `N_samples` is the integrated autocorrelation time. A good guess of the autocorrelation time based on previous simulations for smaller system sizes or nearby parameter points can result in substantial computational saving costs in production runs for large system sizes. Near continuous phase transitions, the integrated autocorrelation time has an additional dependence  $\tau_{\text{int}} \sim L^z$ , where  $z$  is the dynamical exponent of the universality class of the transition. For a 2D system in the vicinity of a continuous phase transition, a sensible scaling for `N_between_samples` could be  $\mathcal{O}(L^2 \times \beta \times L^z)$  ( $\mathcal{O}(L^2)$  links, each has off-diagonal spin flips  $\mathcal{O}(\beta)$ ).

#### 4.3.7 Choosing `N_resamples`

As with `N_samples`, more is better (but also more costly). Usually `N_resamples`  $\approx 1000$  is a sensible choice.

#### 4.3.8 Extracting snapshots

When the option `save_snapshots` is enabled, ParaToric will write the snapshots into the directory specified in `path_out` (or in the paths `paths_out` for hysteresis sweeps). The snapshots are saved in the GraphML format (XML-based), which is supported by many major graph libraries. One snapshot will be saved for every measurement of observables, i.e., `N_samples` snapshots in total. All snapshots are written into a single file to save disk space and simultaneously offer a structured, self-documenting format. Every edge stores a list of spins. The first spin belongs to the first snapshot, the second one to the second snapshot, and so on. There are no special requirements for disks or memory bandwidth; the snapshots are kept in RAM and are only written to disk after the simulation has finished.

#### 4.3.9 Adding new observables/lattices/updates

After adding features to the code, *always* benchmark them using analytical results, other numerical methods (exact diagonalization, tensor networks, ...), and unit tests. We advise using a fixed seed during development, e.g., when checking whether two methods produce the exact same result. The code has some built-in features to check self-consistency, e.g., at the end of each simulation, the code checks whether the cached total energy is numerically close to the total energy calculated from scratch. Do not turn off these features, as they will point you toward bugs!

### 4.4 Benchmarks

#### 4.4.1 Thermalization

In Fig. 2 (which we already discussed before and repeat here for completeness) we plot the gauge field energy  $\propto \lambda$  for two systems: one is sufficiently thermalized, the other one is not. The plots are a direct output of the Python command-line interface. Always make sure that the system is well thermalized.



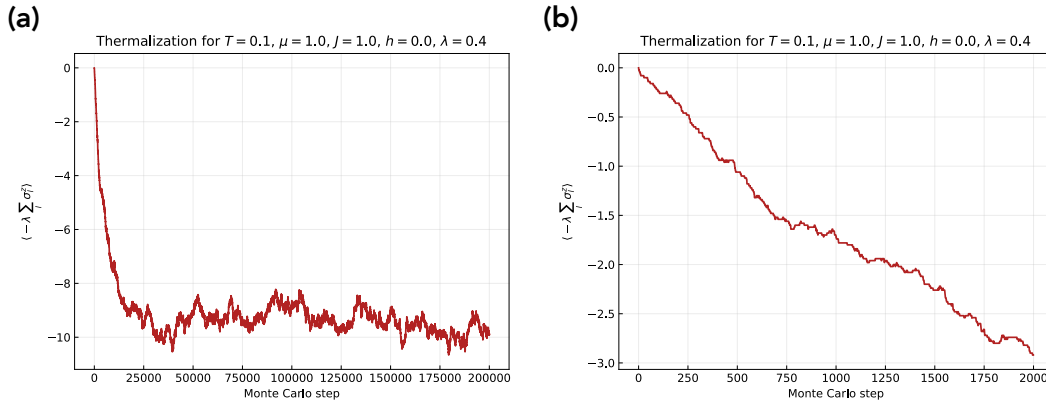


Figure 2: **Good and bad thermalization** plots produced by the Python command-line interface. We show the gauge field energy  $\propto \lambda$ . (a) The system is well thermalized; after its initial drop, the energy fluctuates around the expectation value. (b) The system is not yet thermalized; the floating average is still decreasing.

#### 892 4.4.2 Integrated autocorrelation time

893 Here, we demonstrate how the integrated autocorrelation time  $\tau_{\text{int}}$  grows with decreasing  $N_{\text{between\_samples}}$ . We use the following setup:

Listing 13:  $N_{\text{between\_samples}}$  benchmarking setup

```
895 >>> import numpy as np
896 >>> from paratoric import extended_toric_code as etc
897 >>> series, mean, std, binder, binder_std, tau_int = extended_toric_code.
898   get_sample(N_samples=100000, N_thermalization=10000, N_between_samples
899   =1, N_resamples=1000, custom_therm=False, observables=["energy"], seed
900   =0, mu=1.0, h=0.0, h_therm=0.0, J=1.0, lmbda=0, lmbda_therm=0.0, basis='
901   x', lattice_type="square", system_size=4, beta=10, boundaries="periodic"
902   , default_spin=1, save_snapshots=False)
903
```

905 We only run the simulation once per  $N_{\text{between\_samples}}$ . The results are:

$N_{\text{between\_samples}}$	1	10	100	500	1000
$\tau_{\text{int}}$ (energy)	1895	141.4	19.7	3.24	1.64

907 For very small  $N_{\text{between\_samples}}$ ,  $\tau_{\text{int}}$  is very high: in cases where the update is re-  
 908jected, the configuration is identical to the one measured before! A choice of  $N_{\text{between\_s}}$   
 909amples between 500 and 1000 would be a good tradeoff between  $\tau_{\text{int}}$  and runtime for this  
 910example. Increasing  $N_{\text{between\_samples}}$  to well over 1000 would be a waste of CPU time.

#### 911 4.4.3 Run-time

912 We benchmark the run-time for two realistic parameter sets on the square lattice and varying  
 913system size. The first setup simulates the toric code without fields:<sup>3</sup>

Listing 14:  $L$  benchmarking setup 1

```
914 >>> import numpy as np
915 >>> from paratoric import extended_toric_code as etc
916
```

<sup>3</sup>All tests were run on a laptop; some conditions, like the CPU temperature, were not identical for all simulations. The benchmarks are therefore only an approximation.

```

917 >>> L=20
918 >>> series, mean, std, binder, binder_std, tau_int = extended_toric_code.
919     get_sample(N_samples=10000, N_thermalization=500*L*L*L,
920     N_between_samples=8*L*L*L, N_resamples=1000, custom_therm=False,
921     observables=["energy", "sigma_x", "sigma_z"], seed=0, mu=1.0, h=0.0,
922     h_therm=0.0, J=1.0, lmbda=0, lmbda_therm=0.0, basis='x', lattice_type="
923     square", system_size=L, beta=L, boundaries="periodic", default_spin=1,
924     save_snapshots=False)

```

926 We only run one test per system size. The results are:

$L$	4	8	12	16	20
Runtime (s)	3.1	21.3	75	197	379

928 From our experience, for large  $L$  the update complexity is approximately  $\mathcal{O}(L^3 \log \beta)$ , owing  
929 to the chosen cubic dependency of  $N_{\text{thermalization}}$  and  $N_{\text{between\_samples}}$  and a  
930  $\mathcal{O}(\log \beta)$  dependence of operations on the imaginary time axis, see  $\beta$  benchmark below. The  
931 system size itself does not impact the performance, as the interactions are local. On computing  
932 clusters, we have realized system sizes of up to  $L = 80$  for the square lattice; this number will  
933 only increase in the future as CPUs get faster.

934 The second setup simulates the toric code with fields in both  $\hat{\sigma}^x$  and  $\hat{\sigma}^z$ -direction:

Listing 15:  $L$  benchmarking setup 2

```

935 >>> import numpy as np
936 >>> from paratoric import extended_toric_code as etc
937 >>> L=20
938 >>> series, mean, std, binder, binder_std, tau_int = extended_toric_code.
939     get_sample(N_samples=10000, N_thermalization=500*L*L*L,
940     N_between_samples=8*L*L*L, N_resamples=1000, custom_therm=False,
941     observables=["energy", "sigma_x", "sigma_z"], seed=0, mu=1.0, h=0.2, J
942     =1.0, lmbda=0.2, basis='x', lattice_type="square", system_size=L, beta=L
943     , boundaries="periodic", default_spin=1, save_snapshots=False)
944

```

946 We only run one test per system size. The results are:

$L$	4	8	12	16	20
Runtime (s)	3.9	34.1	133	323	689

948 We also test the run-time dependence of the inverse temperature  $\beta$ , with the following  
949 setup:

Listing 16:  $\beta$  benchmarking setup

```

950 >>> import numpy as np
951 >>> from paratoric import extended_toric_code as etc
952 >>> series, mean, std, binder, binder_std, tau_int = extended_toric_code.
953     get_sample(N_samples=10000, N_thermalization=20000, N_between_samples
954     =2000, N_resamples=1000, custom_therm=False, observables=["energy", "
955     sigma_x", "sigma_z"], seed=0, mu=1.0, h=0.2, J=1.0, lmbda=0.2, basis='x'
956     , lattice_type="square", system_size=10, beta=20, boundaries="periodic",
957     default_spin=1, save_snapshots=False)
958

```

960 We only run one test per  $\beta$ . The results are:

$\beta$	4	8	12	16	20
Runtime (s)	14.9	17.2	19.1	20.0	22.1

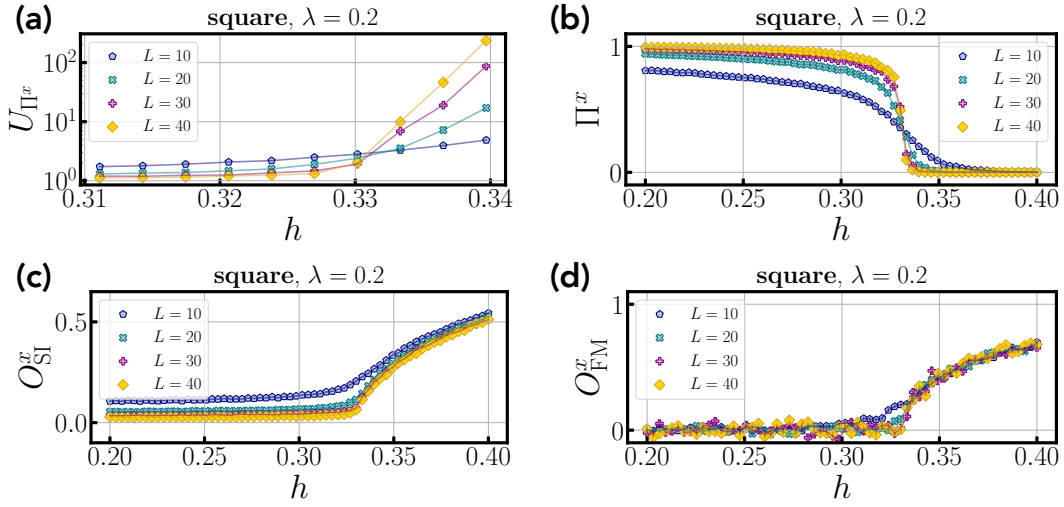


Figure 3: **Topological phase transition** of the extended toric code (1) on the square lattice. The critical field is known and located at  $h_c(\lambda = 0.2) \approx 0.33$  [5, 9]. Our results agree with the value published in the literature, within error bars. (a) The percolation probability Binder ratio  $U_{\Pi^x}$  [9] features a crossing point around  $h = 0.33$ . (b) The percolation probability  $\Pi^x$  is non-zero in the topological phase and zero in the trivial phase. The transition gets sharper with increasing system size. (c) The staggered imaginary time order parameter  $O_{SI}^x$  is zero in the topological phase and non-zero in the trivial phase. (d) The Fredenhagen-Marcu order parameter  $O_{FM}^x$  is zero in the topological phase and non-zero in the trivial phase. The loop length grows with  $\mathcal{O}(L)$ . Compared to the other order parameters, it is very noisy because it is a multi-body correlator and on top of that a division of two exponentially small numbers.

962 This benchmark illustrates an appealing feature of our implementation: there is almost no  
 963 slowing down when increasing  $\beta$  implying that very low temperatures are within reach with  
 964 ParaToric. This paradoxical result (because the number  $n$  of off-diagonal star/plaquette and  
 965 magnetic field operators must physically scale linearly in  $\beta$ ) is explained by the fact that most  
 966 searches within the imaginary time axis scale as  $\mathcal{O}(\log n)$  by making use of binary searches.

#### 967 4.4.4 Topological phase transition

968 We probe the well-known topological phase transition in the ground state of the extended toric  
 969 code (1) on the square lattice, where we have a gapped  $\mathbb{Z}_2$  quantum spin liquid for small fields  
 970  $h, \lambda$  and a topologically trivial phase for high fields. We set  $J = \mu = 1$ ,  $\lambda = 0.2$  and sweep  $h$   
 971 over the known critical value  $h_c(\lambda = 0.2) \approx 0.33$  [5, 9] for  $L \in \{10, 20, 30, 40\}$  in the  $\hat{\sigma}^x$ -basis.  
 972 The temperature is set to  $T = 1/L$  to capture ground state physics. We take 30000 snapshots,  
 973 with  $8L^3$  steps in between snapshots and  $500L^3$  thermalization steps.<sup>4</sup> We confirm that the  
 974 system is well thermalized and that all integrated autocorrelation times are below 10, i.e.,  
 975 the produced snapshots can safely be considered identically and independently distributed.  
 976 We show the percolation probability, the Fredenhagen-Marcu string order parameter, and the  
 977 staggered imaginary time order parameter in Fig. 3. All of them reproduce the known phase  
 978 boundary.

<sup>4</sup>If we were interested in quantities like critical exponents and we need to go very close to the critical field, we should take into account the dynamical exponent  $z$  ( $\tau_{\text{int}} \sim L^z$ ) in the number of steps between snapshots to account for critical slowing down. Now the error bars are just larger near the critical field.

## 5 Conclusion & Outlook

We have presented ParaToric, a continuous-time quantum Monte Carlo solver for the toric code in a parallel field. ParaToric builds on the existing work by Wu, Deng, and Prokof'ev [5] and is also applicable to high temperature and low off-diagonal couplings.

ParaToric can store snapshots, which makes it ideally suited to generate training/benchmarking data for applications in other fields, such as lattice gauge theories, cold atom or other quantum simulators, quantum spin liquids, artificial intelligence, and quantum error correction. We believe it also serves a pedagogical purpose. Another strength of ParaToric is its interoperability with other programming languages. The C interface is compatible with virtually all programming languages, thus ParaToric can be seamlessly integrated into other projects.

ParaToric comes with an MIT license. For future release of ParaToric we plan extensions along the following lines:

- Additional lattices such as the kagome and the ruby lattice. Given the underlying graph structure used in ParaToric, such extensions are straightforward.
- Additional observables: we think here of, for instance, the finite temperature extension of the fidelity susceptibility to diagnose the phase transitions in the absence of a local order parameter. It would be worthwhile to have additional off-diagonal observables such as the off-diagonal Fredenhagen-Marcu string operators, or correlation functions between off-diagonal operators in space and or time. Measurements of the Renyi entropy are also high on the to-do list. The latter two classes require however major changes to the code, and testing.
- Additional interaction types. There are many classes of models in which topological order may be emergent instead of explicit as in the toric code. Such models typically have additional interactions than the ones covered in ParaToric, such as longer-range Ising interactions, and miss some others (typically the plaquette type interactions, and sometimes even the star terms). It is in general an open problem how to efficiently simulate such models at the lowest temperatures (even for sign-free models). Extending ParaToric to dealing with other types of interactions can thus serve as an additional tool for benchmarking purposes and algorithmic exploration.

## Acknowledgements

The authors acknowledge fruitful discussions with A. Bohrdt, G. De Paciani, G. Dünneweber, F. Grusdt, L. Homeier, and N. V. Prokof'ev.

**Author contributions** SML did the main coding and planning work with input from LP. All authors contributed to the writing of the manuscript.

**Funding information** This research was funded by the European Research Council (ERC) under the European Union's Horizon 2020 research and innovation program – ERC Starting Grant SimUcQuam (Grant Agreement No. 948141), and by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under Germany's Excellence Strategy – EXC-2111 – project number 390814868.

## References

- [1] A. Kitaev, *Fault-tolerant quantum computation by anyons*, Annals of Physics **303**(1), 2 (2003), doi:[https://doi.org/10.1016/S0003-4916\(02\)00018-0](https://doi.org/10.1016/S0003-4916(02)00018-0).
- [2] E. Dennis, A. Kitaev, A. Landahl and J. Preskill, *Topological quantum memory*, Journal of Mathematical Physics **43**(9), 4452 (2002), doi:[10.1063/1.1499754](https://doi.org/10.1063/1.1499754), [https://pubs.aip.org/aip/jmp/article-pdf/43/9/4452/19183135/4452\\_1\\_online.pdf](https://pubs.aip.org/aip/jmp/article-pdf/43/9/4452/19183135/4452_1_online.pdf).
- [3] A. G. Fowler, M. Mariantoni, J. M. Martinis and A. N. Cleland, *Surface codes: Towards practical large-scale quantum computation*, Phys. Rev. A **86**, 032324 (2012), doi:[10.1103/PhysRevA.86.032324](https://doi.org/10.1103/PhysRevA.86.032324).
- [4] J. B. Kogut, *An introduction to lattice gauge theory and spin systems*, Rev. Mod. Phys. **51**, 659 (1979), doi:[10.1103/RevModPhys.51.659](https://doi.org/10.1103/RevModPhys.51.659).
- [5] F. Wu, Y. Deng and N. Prokof'ev, *Phase diagram of the toric code model in a parallel magnetic field*, Phys. Rev. B **85**, 195104 (2012), doi:[10.1103/PhysRevB.85.195104](https://doi.org/10.1103/PhysRevB.85.195104).
- [6] K. Fredenhagen and M. Marcu, *Charged states in  $z_2$  gauge theories*, Communications in Mathematical Physics **92**(1), 81 (1983), doi:[10.1007/BF01206315](https://doi.org/10.1007/BF01206315).
- [7] K. Fredenhagen and M. Marcu, *Confinement criterion for qcd with dynamical quarks*, Phys. Rev. Lett. **56**, 223 (1986), doi:[10.1103/PhysRevLett.56.223](https://doi.org/10.1103/PhysRevLett.56.223).
- [8] K. Fredenhagen and M. Marcu, *Dual interpretation of order parameters for lattice gauge theories with matter fields*, Nuclear Physics B - Proceedings Supplements **4**, 352 (1988), doi:[https://doi.org/10.1016/0920-5632\(88\)90124-7](https://doi.org/10.1016/0920-5632(88)90124-7).
- [9] S. M. Linsel, A. Bohrdt, L. Homeier, L. Pollet and F. Grusdt, *Percolation as a confinement order parameter in  $z_2$  lattice gauge theories*, Phys. Rev. B **110**, L241101 (2024), doi:[10.1103/PhysRevB.110.L241101](https://doi.org/10.1103/PhysRevB.110.L241101).
- [10] G. D nnweber, S. M. Linsel, A. Bohrdt and F. Grusdt, *Percolation renormalization group analysis of confinement in  $z_2$  lattice gauge theories*, Phys. Rev. B **111**, 024314 (2025), doi:[10.1103/PhysRevB.111.024314](https://doi.org/10.1103/PhysRevB.111.024314).
- [11] S. M. Linsel, L. Pollet and F. Grusdt, *Independent  $e$ - and  $m$ -anyon confinement in the parallel field toric code on non-square lattices*, doi:[10.48550/arXiv.2504.03512](https://doi.org/10.48550/arXiv.2504.03512) (2025).
- [12] M. Matsumoto and T. Nishimura, *Mersenne twister: a 623-dimensionally equidistributed uniform pseudo-random number generator*, ACM Trans. Model. Comput. Simul. **8**, 3 (1998).
- [13] U. Wolff, *Critical slowing down*, Nuclear Physics B - Proceedings Supplements **17**, 93 (1990), doi:[https://doi.org/10.1016/0920-5632\(90\)90224-I](https://doi.org/10.1016/0920-5632(90)90224-I).
- [14] D. N. Politis and J. P. Romano, *The stationary bootstrap*, Journal of the American Statistical Association **89**(428), 1303 (1994), doi:[10.1080/01621459.1994.10476870](https://doi.org/10.1080/01621459.1994.10476870).
- [15] D. N. Politis and H. White, *Automatic block-length selection for the dependent bootstrap*, Econometric Reviews **23**(1), 53 (2004), doi:[10.1081/ETC-120028836](https://doi.org/10.1081/ETC-120028836).
- [16] A. Patton, D. N. Politis and H. White, *Correction to "automatic block-length selection for the dependent bootstrap" by d. politis and h. white*, Econometric Reviews **28**(4), 372 (2009), doi:[10.1080/07474930802459016](https://doi.org/10.1080/07474930802459016).